

## function approximation

goal: architectures & algorithms  
for approximating (optimal)  
value and policy functions

Bertsekas & Tsitsiklis ch 3

- we'd like to consider MDP

$$\min_u E[c(x,u)] \text{ s.t. } x^+ \sim P(x,u)$$

with infinite state and action  
spaces, e.g.  $X = \mathbb{R}^n$ ,  $U = \mathbb{R}^m$

- now value  $v: X \rightarrow \mathbb{R}$  and

policy  $\mu: X \rightarrow \Delta(U)$

cannot be represented with  
a finite-dimensional vector

\* to use a digital computer,  
must approximate these functions  
i.e. need: - architectures  
- algorithms

# architectures for approximation

- we seek to approximate  $f: X \rightarrow \mathbb{R}$  using  $\tilde{f}: X \times \Theta \rightarrow \mathbb{R}$

(e.g.  $f = v^*, \pi^*, \mu, v^\mu$ )

where  $\theta \in \Theta$  are approximator parameters

→ what properties are desirable?

- provide good approximation of  $f$
- easy to choose parameters  $\theta$ , e.g. linear (convex, at least)

- Broadly, can consider architectures that are linear or nonlinear in parameters  $\theta$

we'll use semicolon to remind that  $\theta$  is a parameter

- linear:  $\tilde{f}(x; \theta) = \sum_{k=0}^K \theta_k \cdot b_k(x)$

where  $\{b_k\}_{k=0}^K$  are called


basis functions

(though they might not form a

vector space basis)

- nonlinear: huge variety of options;  
(so-called) "neural" networks are  
popular:


ex: (single-layer perceptron)



$$\tilde{f}(x; \theta) = \sum_{k=1}^K \theta_k \cdot \sigma \left( \sum_{l=1}^L \theta_{k,l} \cdot z_l(x) \right)$$


where  $\sigma: \mathbb{R} \rightarrow \mathbb{R}$  is sigmoidal,  
i.e. differentiable, monotonic,

$$-\infty < \lim_{\tau \rightarrow -\infty} \sigma(\tau) < \lim_{\tau \rightarrow +\infty} \sigma(\tau) < +\infty$$



$$\frac{e^{\tau} - e^{-\tau}}{e^{\tau} + e^{-\tau}} \quad \text{hyperbolic tangent}$$

$$\frac{1}{1 + e^{-\tau}} \quad \text{logistic}$$



$$\sigma(-\infty) \quad \frac{1}{1+e^{-z}} \quad \text{logistic}$$

and  $z: X \rightarrow \mathbb{R}^L$  is feature vector,  
i.e. (partial) state observation/  
(approximate) sufficient statistic in  
a finite-dimensional vector space

fact: any continuous function over  
a closed and bounded domain can  
be approximated arbitrarily well  
by single-layer perceptron  
(with sufficiently large  $K, L$  and  
the right choice of weights)

\* multiple layers are allowed;  
"in practice, the number of hidden  
layers is usually one or two,  
and almost never more than three" ✓

idea: back propagation (ie chain rule)

• for a multilayer network, function is composition

$$\tilde{f}(x; \theta) = \sigma_N(w_N \cdot \sigma_{N-1}(w_{N-1} \cdots \sigma_1(w_1 \cdot x) \cdots))$$

• so gradient wrt  $\theta$  is product

◦ so gradient wrt  $\Theta$  is product

$$D_{\Theta} \tilde{f} = [D\sigma_N \cdot DW_N \cdot D\sigma_{N-1} \cdot DW_{N-1} \cdots D\sigma_1 \cdot DW_1](x; \Theta)$$

\* forward pass to evaluate derivatives, then backward pass to compute vector-matrix multiplication

### algorithms for approximation

◦ assuming there exists  $\Theta^* \in \Theta$  for which  $\tilde{f}_{\Theta^*} \simeq f$ , how to find  $\Theta^*$ ?

→ propose an algorithmic approach to compute / approximate  $\Theta^*$   
(specify algorithm inputs & any assumptions)

- formulate an optimization problem:

$$\min_{\Theta \in \Theta} \underbrace{\|\tilde{f}_{\Theta} - f\|_{\mathbb{R}^X}^2}_{\text{norm on function space } \mathbb{R}^X}$$

- since norm  $\|\cdot\|_{\mathbb{R}^X}$  can't be evaluated on infinite-dimensional  $\mathbb{R}^X$ , assume given a finite dataset

$$\{(\mathbf{x}, v(\mathbf{x}))\}_{\mathbf{x} \in \Xi} \text{ where } \mathbf{x} \in X$$

$\{(\mathbf{x}, v(\mathbf{x}))\}_{\mathbf{x} \in \mathcal{X}}$  where  $\mathbf{x} \in \mathcal{X}$   
 and  $v(\mathbf{x}) \simeq v^*(\mathbf{x})$  or  $v^\pi(\mathbf{x})$ ,  
 i.e.  $v(\mathbf{x})$  is an estimate of  
 optimal value or value of policy  
 $\pi$  at state  $\mathbf{x} \in \mathcal{X}$ :

$$\min_{\theta \in \Theta} \sum_{\mathbf{x}} \|\tilde{v}(\mathbf{x}; \theta) - v(\mathbf{x})\|^2$$

- can then apply standard algorithms  
 to this finite-dimensional (nonlinear)  
 least-squares optimization problem

ex: linear case (3.2.2)

• suppose  $\tilde{f}(\mathbf{x}; \theta) = \sum_{k=1}^N \theta_k b_k(\mathbf{x})$

$$\Rightarrow D_\theta \tilde{f}(\mathbf{x}; \theta) = [b_1(\mathbf{x}) \cdots b_N(\mathbf{x})]$$

• then  $\min_{\theta} \sum_{\mathbf{x}} \|\tilde{f}(\mathbf{x}; \theta) - f(\mathbf{x})\|_2^2$  is  
 linear least-squares problem

→ solve this NLP

- stationarity:  $\sum_{\mathbf{x}} (\tilde{f}(\mathbf{x}; \theta) - f(\mathbf{x})) D\tilde{f}(\mathbf{x}, \theta)$

- stationarity :  $\sum_x (f(x; \theta) - f(x))' D f(x, \theta)$

$$= \sum_x \left( \underbrace{\sum_{k=1}^N \theta_k b_k(x) - f(x)}_{* \text{ affine in } \theta} \right) [b_1(x) \dots b_N(x)]$$

$$= A \cdot \theta + b = 0 \Leftrightarrow \theta = -A^{-1}b$$